

# **EN2150 - Communication Network Engineering**



**SDN Lab (Individual)**

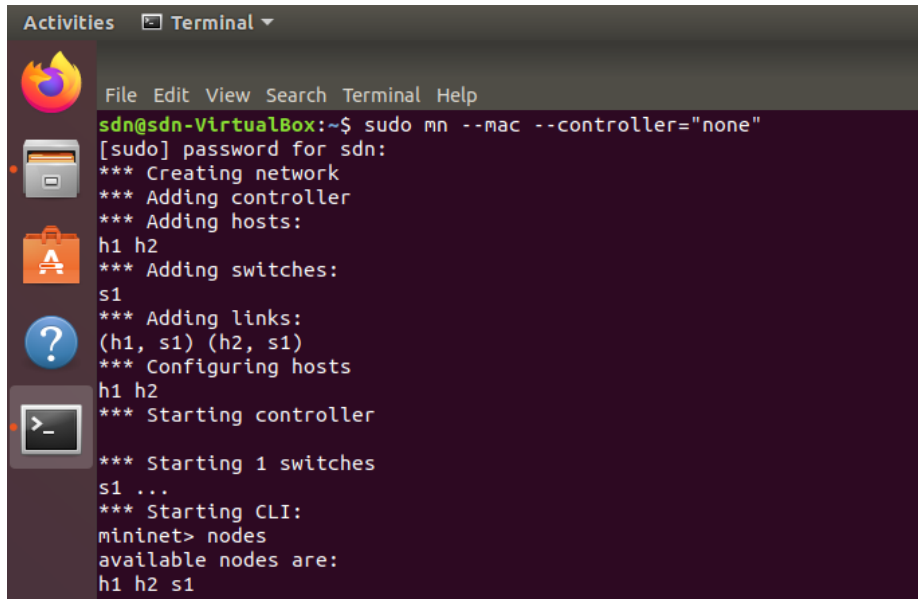
**210129P - Dilshan N.L**

# OpenFlow Configuration Lab

## Part – 01

1. Create a simple network on the Mininet:

`sudo mn --mac --controller="none"`



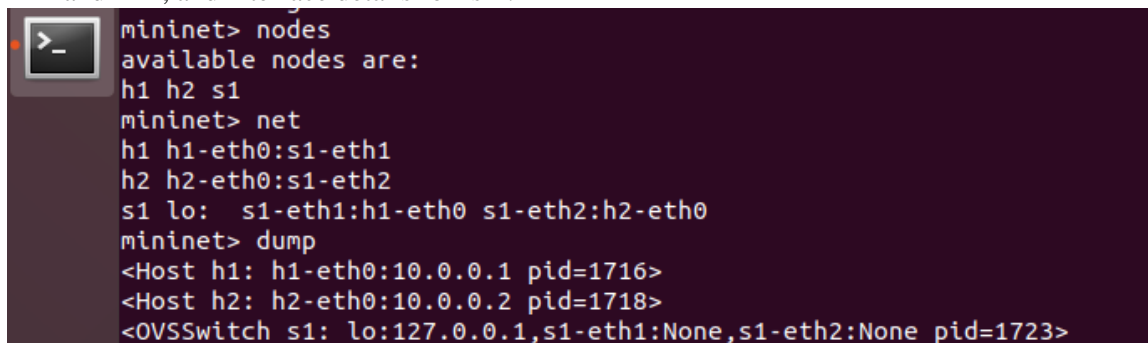
```
sdn@sdn-VirtualBox:~$ sudo mn --mac --controller="none"
[sudo] password for sdn:
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet> nodes
available nodes are:
h1 h2 s1
```

This command starts a Mininet network with two hosts (h1, h2) and one switch (s1) without a controller. It creates and configures the network, and then starts the Mininet CLI.

2. Try the following commands in the Mininet CLI to get an understanding of the network:

```
mininet> nodes
mininet> net
mininet> dump
```

The 'nodes' command lists the nodes: 'h1', 'h2', and 's1'. The 'net' command shows that 'h1' and 'h2' are connected to 's1'. The 'dump' command provides detailed info, including IPs for 'h1' and 'h2', and interface details for 's1'.



```
mininet> nodes
available nodes are:
h1 h2 s1
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=1716>
<Host h2: h2-eth0:10.0.0.2 pid=1718>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=1723>
```

### 3. Attempt pings between the hosts:

```
mininet> h1 ping h2
```

```
mininet> h2 ping h1
```

```
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
From 10.0.0.1 icmp_seq=1 Destination Host Unreachable
From 10.0.0.1 icmp_seq=2 Destination Host Unreachable
From 10.0.0.1 icmp_seq=3 Destination Host Unreachable
From 10.0.0.1 icmp_seq=4 Destination Host Unreachable
From 10.0.0.1 icmp_seq=5 Destination Host Unreachable
From 10.0.0.1 icmp_seq=6 Destination Host Unreachable
^C
--- 10.0.0.2 ping statistics ---
7 packets transmitted, 0 received, +6 errors, 100% packet loss, time 6142ms
pipe 4
mininet> h2 ping h1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
From 10.0.0.2 icmp_seq=1 Destination Host Unreachable
From 10.0.0.2 icmp_seq=2 Destination Host Unreachable
From 10.0.0.2 icmp_seq=3 Destination Host Unreachable
From 10.0.0.2 icmp_seq=4 Destination Host Unreachable
From 10.0.0.2 icmp_seq=5 Destination Host Unreachable
From 10.0.0.2 icmp_seq=6 Destination Host Unreachable
^C
--- 10.0.0.1 ping statistics ---
7 packets transmitted, 0 received, +6 errors, 100% packet loss, time 6144ms
pipe 4
```

The ping attempts failed because the network setup does not include an SDN controller. In an SDN environment, the controller is essential for managing packet flow between hosts. Since the network was initialized with the `--controller="none"` option, the switch lacks the necessary flow rules to route packets between the hosts. As a result, the absence of a controller prevented the successful transmission of pings between the hosts.

### 4. Check the flow table of the switch:

```
mininet> dpctl dump-flows
```

```
pipe 4
mininet> dpctl dump-flows
*** s1 ***
mininet>
```

The flow table is empty because there's no SDN controller connected to the switch to add rules. The SDN controller normally populates the flow table with rules to direct how packets should be handled. Without these rules, the switch doesn't know how to process incoming packets, which is why you're seeing "Destination Host Unreachable" messages during ping attempts. Essentially, the lack of flow rules means the switch isn't set up to handle traffic, causing the ping failures.

### 5. Exit from the Mininet CLI:

```
mininet> exit
```

```
cookie=0x0, duration=250.000s, table=0, n_packets=50, n_bytes=2700, priority=0 actions=CONTROLLER
mininet> exit
*** Stopping 1 controllers
c0
*** Stopping 2 links
..
*** Stopping 1 switches
s1
*** Stopping 2 hosts
h1 h2
*** Done
completed in 393.504 seconds
sdnnode@fluxus100:~$
```

6. Create another simple network on the Mininet VM but this time use the default controller:

`sudo mn --mac`

```
sdn@sdn-VirtualBox:~$ sudo kill -9 915
sdn@sdn-VirtualBox:~$ sudo ss -ltnp | grep 6653
sdn@sdn-VirtualBox:~$ sudo mn --mac
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2
*** Adding switches:
s1
*** Adding links:
(h1, s1) (h2, s1)
*** Configuring hosts
h1 h2
*** Starting controller
c0
*** Starting 1 switches
s1 ...
*** Starting CLI:
mininet>
```

7. Try the following commands in the Mininet CLI to get an understanding of the network:

`mininet> nodes`

`mininet> net`

`mininet> dump`

```
mininet> net
h1 h1-eth0:s1-eth1
h2 h2-eth0:s1-eth2
s1 lo: s1-eth1:h1-eth0 s1-eth2:h2-eth0
mininet> dump
<Host h1: h1-eth0:10.0.0.1 pid=1716>
<Host h2: h2-eth0:10.0.0.2 pid=1718>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None pid=1723>
```

8. Check the flow table of the switch:

`mininet> dpctl dump-flows`

```
c0
mininet> dpctl dump-flows
*** s1 ***
cookie=0x0, duration=25.326s, table=0, n_packets=14, n_bytes=1172, priority=0 actions=CONTROLLER:128
mininet>
```

The output shows a single flow entry in the switch's flow table. Here's a breakdown of the details:

- **Cookie:** The unique identifier for this flow entry is 0x0.
- **Duration:** The flow entry has been active for 25.326 seconds.
- **Table:** This entry is in table number 0.
- **Packets and Bytes:** The entry has matched 14 packets and 1172 bytes of data.
- **Priority:** The priority level of the flow entry is 0.
- **Actions:** The action for this flow is **CONTROLLER:128**, meaning packets matching this flow are sent to the SDN controller for further processing.
- This setup indicates that the switch is configured to forward certain packets to the SDN controller instead of handling them locally. This often happens when there are specific rules or default actions that direct packets to the controller for more detailed handling or decision-making.

## 9. Attempt pings between the hosts

```
mininet> h1 ping h2
mininet> h2 ping h1
```

The pings between hosts h1 and h2 were successful, as shown by the results. This success indicates that network communication between the two hosts is working properly. The successful pings are due to the presence of an SDN controller in the network setup. The controller manages the flow rules, ensuring that packets are correctly routed between hosts.

In Software-Defined Networking (SDN), the controller is crucial because it determines how switches handle and forward packets. Without the controller, the switch would not have the necessary rules to route packets between hosts, leading to failed pings. With the controller in place, the switch can dynamically set up its flow table to enable communication between h1 and h2.

The ping results show that the first ping from h2 to h1 took slightly longer than subsequent pings. This delay is mainly due to the Address Resolution Protocol (ARP) process:

- **ARP Request:** When the first ping is sent, h2 needs to find out the MAC address of h1. It does this by broadcasting an ARP request to the network asking for the MAC address corresponding to 10.0.0.1.
- **ARP Reply:** h1 responds to the ARP request with its MAC address. Once h2 receives this information, it can send the ping packet to h1.
- **Subsequent Pings:** After the initial ARP process, h2 caches the MAC address of h1. Therefore, later pings do not require ARP resolution, leading to quicker response times and lower latency for subsequent packets.

Overall, the network communication is functioning correctly, and the presence of the SDN controller allows for effective packet forwarding and minimal delay after the initial ARP resolution.

```

mininet> dpctl dump-flows
*** s1 -----
cookie=0x0, duration=25.326s, table=0, n_packets=14, n_bytes=1172, priority=0 actions=CONTROLLER:128
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=2.03 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.244 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.071 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.068 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.076 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.070 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.072 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=0.078 ms
64 bytes from 10.0.0.2: icmp_seq=9 ttl=64 time=0.077 ms
64 bytes from 10.0.0.2: icmp_seq=10 ttl=64 time=0.072 ms
^C
--- 10.0.0.2 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9189ms
rtt min/avg/max/mdev = 0.068/0.286/2.035/0.585 ms
mininet> h2 ping h1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=1.98 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=0.074 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=0.072 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=64 time=0.070 ms
64 bytes from 10.0.0.1: icmp_seq=5 ttl=64 time=0.084 ms
64 bytes from 10.0.0.1: icmp_seq=6 ttl=64 time=0.082 ms
64 bytes from 10.0.0.1: icmp_seq=7 ttl=64 time=0.092 ms
64 bytes from 10.0.0.1: icmp_seq=8 ttl=64 time=0.081 ms
64 bytes from 10.0.0.1: icmp_seq=9 ttl=64 time=0.082 ms
^C
--- 10.0.0.1 ping statistics ---
9 packets transmitted, 9 received, 0% packet loss, time 8152ms
rtt min/avg/max/mdev = 0.070/0.291/1.983/0.598 ms
mininet>

```

## 10. Check the flow table of the switch again:

mininet> dpctl dump-flows

```

mininet> dpctl dump-flows
*** s1 -----
cookie=0x0, duration=38.411s, table=0, n_packets=0, n_bytes=0, idle_timeout=60, priority=1,arp,in_port="s1-eth2",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,arp_spa=10.0.0.2,arp_tpa=10.0.0.1,arp_op=2 actions=output:"s1-eth1"
cookie=0x0, duration=33.222s, table=0, n_packets=0, n_bytes=0, idle_timeout=60, priority=1,arp,in_port="s1-eth2",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,arp_spa=10.0.0.2,arp_tpa=10.0.0.1,arp_op=1 actions=output:"s1-eth1"
cookie=0x0, duration=33.226s, table=0, n_packets=0, n_bytes=0, idle_timeout=60, priority=1,arp,in_port="s1-eth1",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02,arp_spa=10.0.0.1,arp_tpa=10.0.0.2,arp_op=2 actions=output:"s1-eth2"
cookie=0x0, duration=38.411s, table=0, n_packets=9, n_bytes=882, idle_timeout=60, priority=1,icmp,in_port="s1-eth1",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02,nw_src=10.0.0.2,nw_dst=10.0.0.1,nw_tos=0,icmp_type=8,icmp_code=0 actions=output:"s1-eth2"
cookie=0x0, duration=38.410s, table=0, n_packets=9, n_bytes=882, idle_timeout=60, priority=1,icmp,in_port="s1-eth2",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,nw_src=10.0.0.2,nw_dst=10.0.0.1,nw_tos=0,icmp_type=0,icmp_code=0 actions=output:"s1-eth1"
cookie=0x0, duration=23.878s, table=0, n_packets=8, n_bytes=784, idle_timeout=60, priority=1,icmp,in_port="s1-eth2",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:02,dl_dst=00:00:00:00:00:01,nw_src=10.0.0.2,nw_dst=10.0.0.1,nw_tos=0,icmp_type=0,icmp_code=0 actions=output:"s1-eth1"
cookie=0x0, duration=23.869s, table=0, n_packets=8, n_bytes=784, idle_timeout=60, priority=1,icmp,in_port="s1-eth1",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:00:01,dl_dst=00:00:00:00:00:02,nw_src=10.0.0.1,nw_dst=10.0.0.2,nw_tos=0,icmp_type=0,icmp_code=0 actions=output:"s1-eth2"
***
cookie=0x0, duration=80.106s, table=0, n_packets=26, n_bytes=2012, priority=0 actions=CONTROLLER:128
mininet>

```

The flow table entries handle ARP and ICMP traffic between hosts. For ARP packets, fields like in port, dl src (source MAC), dl dst (destination MAC), arp spa (source IP), arp tpa (target IP), and arp op (operation) define how ARP requests and replies are processed. For ICMP traffic, fields like nw src (source IP), nw dst (destination IP), icmp type, and icmp code specify the ICMP messages and actions for packet handling.

- Duration: How long the entry has been active.
- n packets and n bytes: Number of packets and bytes matching the flow.
- Idle Timeout: Time before the entry is removed if not used.
- Priority: Determines the order of flow entry matching.
- Actions: Specifies how packets are processed, such as forwarding to specific ports.
- These rules were added as the switch processed ARP and ICMP traffic, allowing it to forward packets correctly and achieve successful pings between the hosts.

- 11. Wait for the flows to timeout. That is, wait until the output of the following command is again empty:**

```
mininet> dpctl dump-flows
```

```
cookie=0x0, duration=259.105s, table=0, n_packets=35, n_bytes=2698, priority=0 actions=CONTROLLER:128
*** s1 ***
mininet> dpctl dump-flows
*** s1 ***
```

- 12. We will now enable snooping to see the message flow between the switch and controller:**

```
mininet> dpctl snoop &
```

```
mininet> dpctl snoop &
*** s1 ***
mininet> dpctl snoop &
*** s1 ***
mininet> █
```

- 13. Wait for the flows to timeout. That is, wait until the output of the following command is again empty:**

```
mininet> dpctl dump-flows
```

- 14. Check the flow table again:**

```
mininet> dpctl dump-flows
```

```
mininet> dpctl dump-flows
*** s1 ***
cookie=0x0, duration=349.542s, table=0, n_packets=30, n_bytes=2292, priority=0 actions=CONTROLLER:128
```

The flow table is empty for a few possible reasons:

- **Idle Timeout:** Flow entries might be removed automatically if no packets match them within a set period. If the entries had an idle timeout and no packets matched, they would be deleted.
- **Table Clearance:** The flow table may have been cleared by the switch or controller, either manually or automatically, to make space for new entries or after a reset.
- **Switch or Controller Actions:** The switch or controller might have cleared the entries as part of its normal operation or management.

Without entries, the switch will handle new packets based on default settings or wait for new instructions from the controller. New flow entries will appear when new traffic is processed, depending on the switch's and controller's setup.

- 15. Pings between the hosts again:**

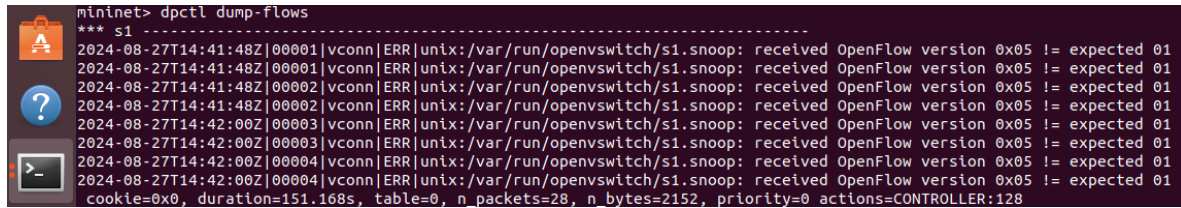
```
mininet> h1 ping h2
```



```
mininet> h2 ping h1
```

## 16. Check the flow table npw:

```
mininet> dpctl dump-flows
```



```
mininet> dpctl dump-flows
*** s1 ***
2024-08-27T14:41:48Z|00001|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:41:48Z|00001|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:41:48Z|00002|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:41:48Z|00002|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:42:00Z|00003|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:42:00Z|00003|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:42:00Z|00004|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
2024-08-27T14:42:00Z|00004|vconn|ERR|unix:/var/run/openvswitch/s1.snoop: received OpenFlow version 0x05 != expected 01
cookie=0x0, duration=151.168s, table=0, n packets=28, n bytes=2152, priority=0 actions=CONTROLLER:128
```

The flow entries in the table were related to specific network packets matched and managed according to rules set by the OpenFlow controller. These entries included attributes such as priority, idle timeout, hard timeout, and actions for handling packets. The duration field showed how long each entry had been active.

As the experiment continued, snooping was enabled and the system waited for the flow entries to time out. Flow entries in OpenFlow switches have an idle timeout, meaning they are removed if no packets match them within a certain period. After the timeout period, running `dpctl dump-flows` showed an empty flow table, indicating that all previous entries had expired or been removed.

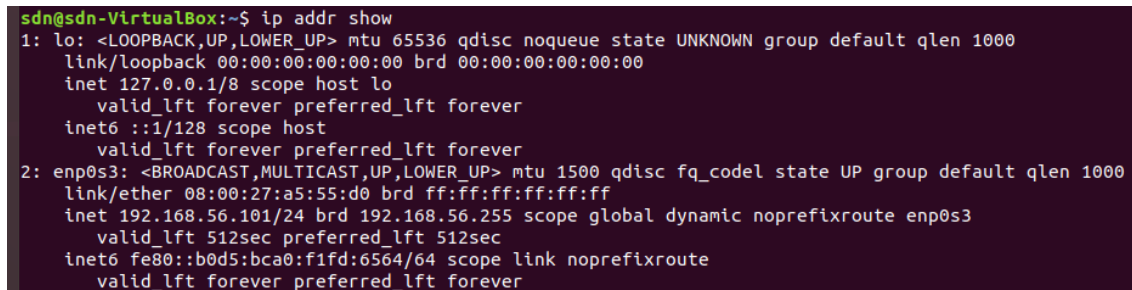
When pings were initiated between the hosts again (`h1 ping h2` and `h2 ping h1`), new flow entries were created in the table by the OpenFlow switch to handle these packets. This is expected, as the switch dynamically adds new entries for incoming packets that don't match existing flows.

This behavior reflects how OpenFlow and SDN manage flow tables: they start with entries based on network activity, clear them after timeouts, and re-populate the table with new entries as needed.

## Part – 02

### 1. Check the IP address configuration of the Ubuntu VM from :

```
ip addr show
```



```
sdn@sdn-VirtualBox:~$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:a5:55:d0 brd ff:ff:ff:ff:ff:ff
    inet 192.168.56.101/24 brd 192.168.56.255 scope global dynamic noprefixroute enp0s3
        valid_lft 512sec preferred_lft 512sec
    inet6 fe80::b0d5:bca0:f1fd:6564/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```



## 2. Open a New terminal window in Ubuntu VM and Run OpenDaylight:

```
sudo ./karaf-0.7.2/bin/karaf
```

```

at java.lang.ProcessImpl.(ProcessImpl.java:65)
at java.lang.ProcessImpl.(ProcessImpl.java:50)
at java.lang.ProcessImpl.(ProcessImpl.java:24)
at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
at java.lang.Class.newInstance(Class.java:442)
at jline.TerminalFactory.getFlavor(TerminalFactory.java:211)
at jline.TerminalFactory.create(TerminalFactory.java:102)
at jline.TerminalFactory.create(TerminalFactory.java:51)
at org.apache.karaf.shell.impl.console.TerminalFactory.init(TerminalFactory.java:37)
at org.apache.karaf.shell.impl.console.TerminalFactory.getTerminal(TerminalFactory.java:30)
at org.apache.karaf.shell.impl.console.osgi.LocalConsoleManager.start(LocalConsoleManager.java:61)
at org.apache.karaf.shell.impl.console.osgi.Activator.start(Activator.java:112)
at org.eclipse.osgi.internal.framework.BundleContextImpl$3.run(BundleContextImpl.java:771)
at org.eclipse.osgi.internal.framework.BundleContextImpl$3.run(BundleContextImpl.java:1)
at java.security.AccessController.doPrivileged(Native Method)
at org.eclipse.osgi.internal.framework.BundleContextImpl.startActivator(BundleContextImpl.java:764)
at org.eclipse.osgi.internal.framework.BundleContextImpl.start(BundleContextImpl.java:721)
at org.eclipse.osgi.internal.framework.EquinoxBundle.startWorker0(EquinoxBundle.java:941)
at org.eclipse.osgi.internal.framework.EquinoxBundle$EquinoxModule.startWorker(EquinoxBundle.java:318)
at org.eclipse.osgi.container.Module.doStart(Module.java:571)
at org.eclipse.osgi.container.Module.start(Module.java:439)
at org.eclipse.osgi.container.ModuleContainer$ContainerStartLevel.incStartLevel(ModuleContainer.java:1582)
at org.eclipse.osgi.container.ModuleContainer$ContainerStartLevel.incStartLevel(ModuleContainer.java:1562)
at org.eclipse.osgi.container.ModuleContainer$ContainerStartLevel.doContainerStartLevel(ModuleContainer.java:1533)
at org.eclipse.osgi.container.ModuleContainer$ContainerStartLevel.dispatchEvent(ModuleContainer.java:1476)
at org.eclipse.osgi.container.ModuleContainer$ContainerStartLevel.dispatchEvent(ModuleContainer.java:1)
at org.eclipse.osgi.framework.eventmgr.EventManager.dispatchEvent(EventManager.java:230)
at org.eclipse.osgi.framework.eventmgr.EventManager$EventThread.run(EventManager.java:340)

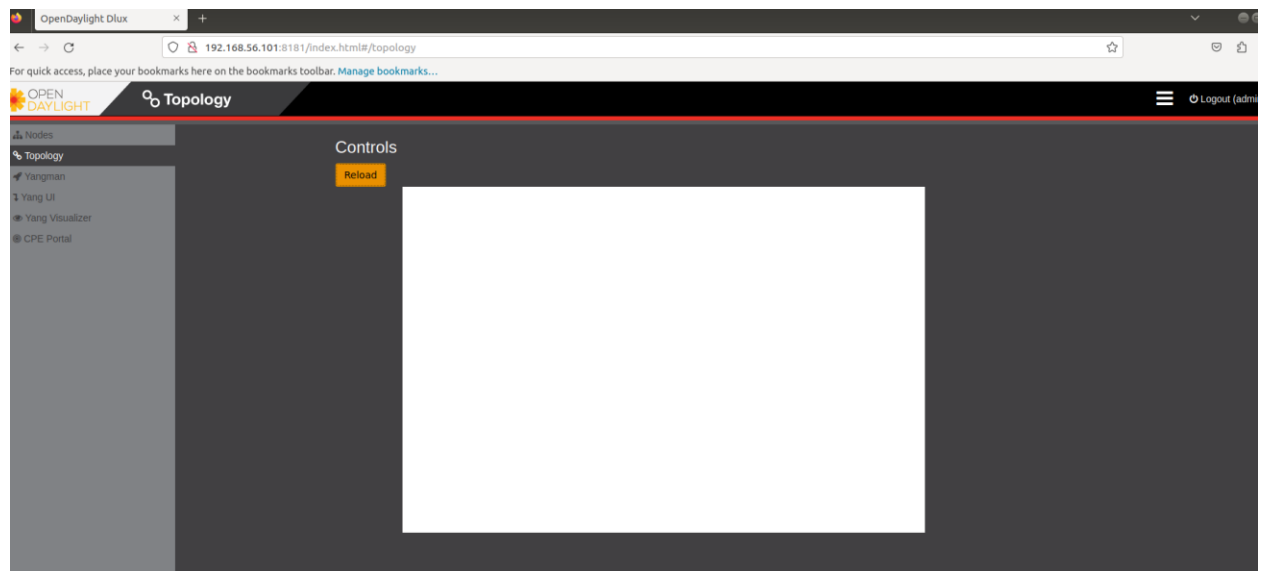
```

```
00% [=====]
araf started in 27s. Bundle stats: 493 active, 494 total
```

[illegible]

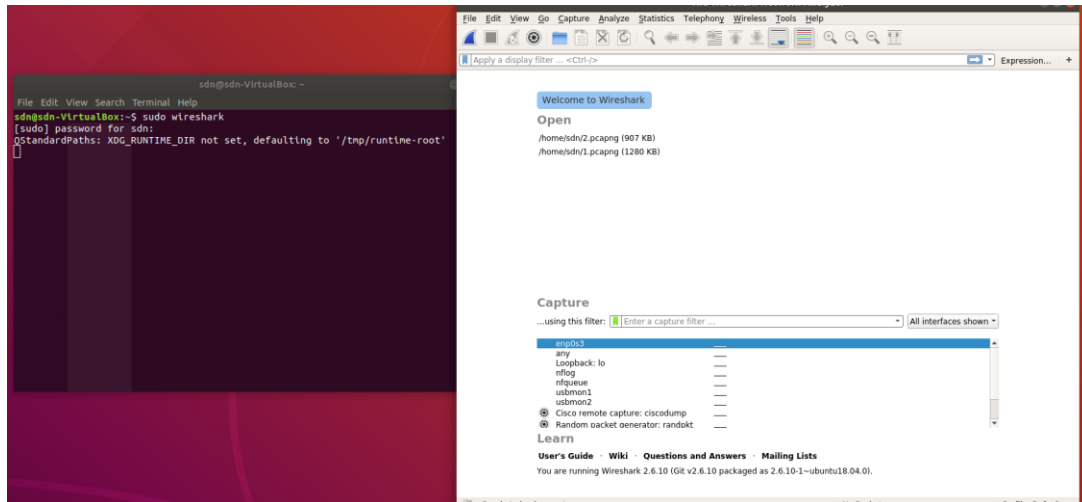
```
It '<tab>' for a list of available commands
nd '[cmd] --help' for help on a specific command.
It '<ctrl-d>' or type 'system:shutdown' or 'logout' to shutdown OpenDaylight.
```

```
pendaylight-user@root>
```

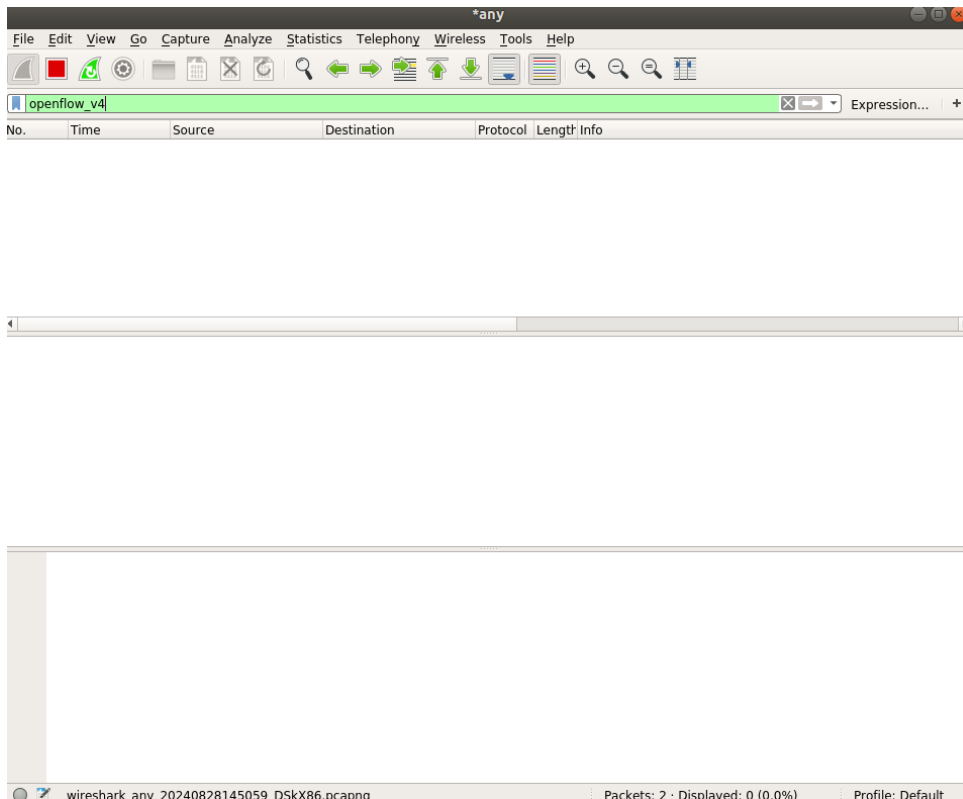


## Part – 03

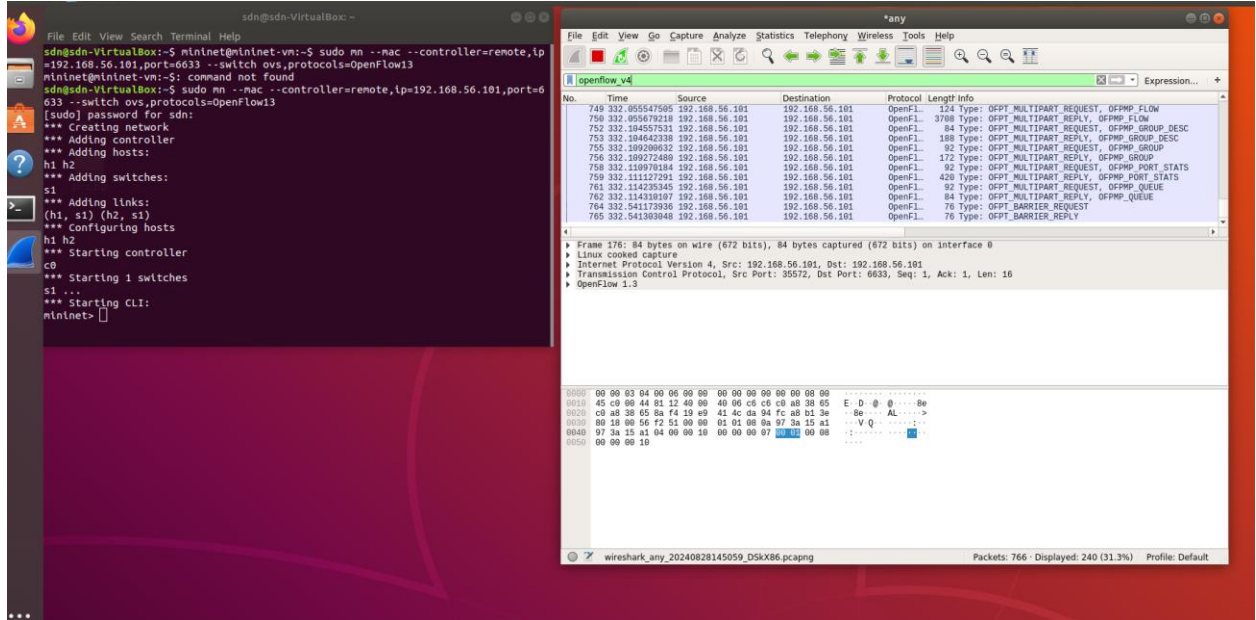
1. Start Wireshark in the Ubuntu VM and select interface corresponding to your LAN:



2. Create a display filter within Wireshark for `openflow_v4`:



### 3. Create a simple network on the Mininet VM:



### 4. Check the flow table of the switch:

mininet> dpctl dump-flows --protocols=OpenFlow13

```
*** Starting CLI:
mininet> dpctl dump-flows --protocols=OpenFlow13
*** s1 ***
cookie=0x2b00000000000000, duration=103.357s, table=0, n_packets=0, n_bytes=0, priority=100,dl_type=0x88cc actions=CONTROLLER:65535
cookie=0x2b00000000000000, duration=101.945s, table=0, n_packets=4, n_bytes=280, priority=2,in_port="s1-eth1" actions=output:"s1-eth2",CONTROLLER:65535
cookie=0x2b00000000000000, duration=101.945s, table=0, n_packets=4, n_bytes=280, priority=2,in_port="s1-eth2" actions=output:"s1-eth1",CONTROLLER:65535
cookie=0x2b00000000000000, duration=103.357s, table=0, n_packets=2, n_bytes=180, priority=0 actions=drop
cookie=0x80800000, duration=103.356s, table=17, n_packets=0, n_bytes=0, priority=0,metadata=0x0000000000000000/0xf000000000000000 actions=write_metadata:0x9000000000000000/0xf000000000000000,goto_table=17
cookie=0x80800000, duration=103.356s, table=18, n_packets=0, n_bytes=0, priority=0 actions=goto_table=38
cookie=0x8220016, duration=103.252s, table=19, n_packets=0, n_bytes=0, priority=100,arp,arp_op=2 actions=resubmit(,17)
cookie=0x8220015, duration=103.252s, table=19, n_packets=0, n_bytes=0, priority=100,arp,arp_op=1 actions=resubmit(,17)
cookie=0x1080000, duration=103.252s, table=19, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x1030000, duration=103.252s, table=20, n_packets=0, n_bytes=0, priority=0 actions=goto_table=80
cookie=0x80800004, duration=103.252s, table=22, n_packets=0, n_bytes=0, priority=0 actions=CONTROLLER:65535
cookie=0x1080000, duration=103.252s, table=23, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x822002d, duration=102.682s, table=43, n_packets=0, n_bytes=0, priority=100,arp,arp_op=1 actions=group:5000
cookie=0x822002e, duration=102.682s, table=43, n_packets=0, n_bytes=0, priority=100,arp,arp_op=2 actions=CONTROLLER:65535,resubmit(,48)
cookie=0x8220000, duration=103.110s, table=45, n_packets=0, n_bytes=0, priority=0 actions=goto_table=48
cookie=0x4000000, duration=103.356s, table=45, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x8500000, duration=103.356s, table=48, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,49),resubmit(,50)
cookie=0x8050001, duration=103.112s, table=50, n_packets=0, n_bytes=0, priority=10,reg4=0x1 actions=goto_table:51
cookie=0x8050000, duration=103.356s, table=50, n_packets=0, n_bytes=0, priority=0 actions=CONTROLLER:65535,learn(table=49,hard_timeout=10,priority=0,cookie=0x800000,NXM_OF_ETH_SRC[],NXM_NX_REG1[0..19],oad:0x1->NXM_NX_REG4[0..7]),goto_table:51
cookie=0x8030000, duration=103.069s, table=51, n_packets=0, n_bytes=0, priority=15,dl_dst=01:80:c2:00:00:00/ff:ff:ff:ff:ff:f0 actions=drop
cookie=0x8030000, duration=103.110s, table=51, n_packets=0, n_bytes=0, priority=0 actions=goto_table:52
cookie=0x8080000, duration=103.356s, table=60, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x1030000, duration=103.252s, table=80, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x8220000, duration=103.078s, table=81, n_packets=0, n_bytes=0, priority=0 actions=drop
cookie=0x4000001, duration=103.357s, table=90, n_packets=0, n_bytes=0, priority=0 actions=resubmit(,17)
cookie=0x6900000, duration=102.985s, table=210, n_packets=0, n_bytes=0, priority=61009,arp actions=drop
cookie=0x6900000, duration=102.970s, table=210, n_packets=0, n_bytes=0, priority=61009,ip actions=drop
cookie=0x6900000, duration=103.356s, table=210, n_packets=0, n_bytes=0, priority=0 actions=drop
cookie=0x6900000, duration=103.356s, table=211, n_packets=0, n_bytes=0, priority=0 actions=goto_table:212
cookie=0x6900000, duration=103.356s, table=212, n_packets=0, n_bytes=0, priority=0 actions=goto_table:213
cookie=0x6900000, duration=103.080s, table=213, n_packets=0, n_bytes=0, priority=62020,ct_state=newest-rel-lnv+trk,ct_mark=0x1/0x1 actions=resubmit(,17)
cookie=0x6900000, duration=103.057s, table=213, n_packets=0, n_bytes=0, priority=62020,ct_state=new-est-rel-lnv+trk,ct_mark=0x1/0x1 actions=resubmit(,17)
cookie=0x6900000, duration=103.356s, table=213, n_packets=0, n_bytes=0, priority=0 actions=drop
cookie=0x6900000, duration=102.935s, table=239, n_packets=0, n_bytes=0, priority=62020,ct_state=trk,ip actions=ct(table=240)
cookie=0x6900000, duration=102.928s, table=239, n_packets=0, n_bytes=0, priority=62020,ct_state=trk,ip actions=ct(table=240)
cookie=0x6900000, duration=102.928s, table=239, n_packets=0, n_bytes=0, priority=61010 actions=goto_table:240
cookie=0x6900000, duration=103.035s, table=240, n_packets=0, n_bytes=0, priority=61010,dl_dst=ff:ff:ff:ff:ff:ff,nw_dst=255.255.255.255 actions=goto_table:241
cookie=0x6900000, duration=103.035s, table=240, n_packets=0, n_bytes=0, priority=01085,dl_dst=ff:ff:ff:ff:ff:ff actions=resubmit(,220)
cookie=0x6900000, duration=103.035s, table=240, n_packets=0, n_bytes=0, priority=0 actions=drop
```

#### 1. LLDP Packets:

- Entry: Handles LLDP packets (dl\_type=0x88cc) used for network topology discovery.
- Priority: 100 (high).
- Action: Forward to the controller.

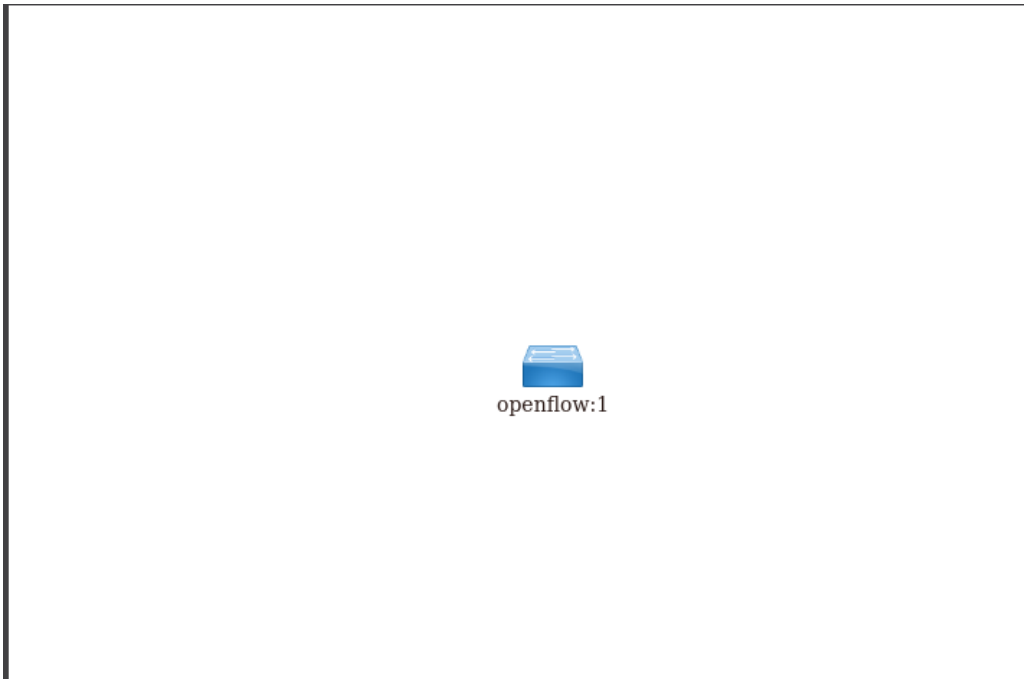
#### 2. Host-to-Host Traffic:

- Entry 1: Matches packets from port s1-eth1 (host h1) and forwards them to s1-eth2 (host h2) and the controller.
  - Entry 2: Matches packets from port s1-eth2 (host h2) and forwards them to s1-eth1 and the controller.
3. Default Drop Rule:
- Entry: Lowest priority (0).
  - Action: Drops any packets that do not match higher-priority rules.

This setup handles LLDP for network discovery, manages traffic between two hosts, and discards unmatched packets.

**5. Open up OpenDaylight GUI:**

<http://192.168.56.101:8181/index.html>



**6. Wireshark FLOW\_MOD:**

Wireshark interface showing a packet capture named 'openflow\_v4'. The packet list displays several OpenFlow messages between 192.168.56.101 and 192.168.56.101. The selected packet (No. 1572) is an OFPT\_MULTIPART\_REQUEST (Type: OFPMP\_TABLE) with a length of 82 bytes. The packet details pane shows the frame structure: Ethernet II, Internet Protocol Version 4, and Transmission Control Protocol. The packet bytes pane displays the raw data in hexadecimal and ASCII.

| No.  | Time          | Source         | Destination    | Protocol  | Length | Info   |
|------|---------------|----------------|----------------|-----------|--------|--|
| 1555 | 178.708327889 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 170    | Type: OFPT_MULTIPART_REPLY, OFPMP_GROUP        |
| 1557 | 178.709198664 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 90     | Type: OFPT_MULTIPART_REQUEST, OFPMP_PORT_STATS |
| 1558 | 178.709322469 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 418    | Type: OFPT_MULTIPART_REPLY, OFPMP_PORT_STATS   |
| 1560 | 178.709932274 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 90     | Type: OFPT_MULTIPART_REQUEST, OFPMP_QUEUE      |
| 1561 | 178.709960162 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 82     | Type: OFPT_MULTIPART_REPLY, OFPMP_QUEUE        |
| 1563 | 179.134998270 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 74     | Type: OFPT_BARRIER_REQUEST                     |
| 1564 | 179.135079394 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 74     | Type: OFPT_BARRIER_REPLY                       |
| 1566 | 180.401788628 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 316    | Type: OFPT_PACKET_OUT                          |
| 1568 | 180.902453447 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 74     | Type: OFPT_BARRIER_REQUEST                     |
| 1570 | 180.902564020 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 74     | Type: OFPT_BARRIER_REPLY                       |
| 1572 | 181.711889636 | 192.168.56.101 | 192.168.56.101 | OpenFl... | 82     | Type: OFPT_MULTIPART_REQUEST, OFPMP_TABLE      |

Frame 1: 82 bytes on wire (656 bits), 82 bytes captured (656 bits) on interface 0  
 Ethernet II, Src: 00:00:00:00:00:00 (00:00:00:00:00:00), Dst: 00:00:00:00:00:00 (00:00:00:00:00:00)  
 Internet Protocol Version 4, Src: 192.168.56.101, Dst: 192.168.56.101  
 Transmission Control Protocol, Src Port: 6633, Dst Port: 57852, Seq: 1, Ack: 1, Len: 16  
 OpenFlow 1.3

```

0000  00 00 00 00 00 00 00 00 00 00 00 00 08 00 45 00  .....E
0010  00 44 46 b7 40 00 40 06 01 e2 c0 a8 38 65 c0 a8  -DF-@-@-...8e-
0020  38 65 19 e9 e1 fc a6 8e 85 e9 06 78 fe f7 80 18  8e-...-x-...
0030  02 00 f2 51 00 00 01 01 08 0a 62 d9 09 75 62 d8  -Q-...-b-ub-
0040  ff 9b 0a 12 00 10 00 00 0c 0b 00 03 00 00 00 00  -.-...-
0050  00 00
  
```

7. Re-start the Wireshark capture:

8. Attempt pinging between h1 and h2:

```

mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=8.58 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.219 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.043 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.064 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.052 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.063 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.079 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=0.047 ms
64 bytes from 10.0.0.2: icmp_seq=9 ttl=64 time=0.042 ms
64 bytes from 10.0.0.2: icmp_seq=10 ttl=64 time=0.045 ms
64 bytes from 10.0.0.2: icmp_seq=11 ttl=64 time=0.045 ms
64 bytes from 10.0.0.2: icmp_seq=12 ttl=64 time=0.038 ms
64 bytes from 10.0.0.2: icmp_seq=13 ttl=64 time=0.040 ms
64 bytes from 10.0.0.2: icmp_seq=14 ttl=64 time=0.274 ms
64 bytes from 10.0.0.2: icmp_seq=15 ttl=64 time=0.072 ms
64 bytes from 10.0.0.2: icmp_seq=16 ttl=64 time=0.068 ms
64 bytes from 10.0.0.2: icmp_seq=17 ttl=64 time=0.045 ms
64 bytes from 10.0.0.2: icmp_seq=18 ttl=64 time=0.046 ms
^C
--- 10.0.0.2 ping statistics ---
18 packets transmitted, 18 received, 0% packet loss, time 17377ms
rtt min/avg/max/mdev = 0.038/0.548/8.584/1.950 ms
mininet> h2 ping h1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=5.04 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=0.058 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=0.048 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=64 time=0.064 ms
64 bytes from 10.0.0.1: icmp_seq=5 ttl=64 time=0.054 ms
^C
--- 10.0.0.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4071ms
rtt min/avg/max/mdev = 0.048/1.054/5.049/1.997 ms
mininet>

```

The successful ping test between h1 and h2 indicates that the OpenFlow switch has been correctly configured. The flow table entries in the switch, as detailed in the previous flow table dump, played a crucial role in this:

Traffic from h1 to h2: A specific flow entry managed traffic from s1-eth1 (connected to h1) to s1-eth2 (connected to h2). This entry ensured that packets were properly forwarded from h1 to h2.

Traffic from h2 to h1: Another flow entry handled traffic from s1-eth2 back to s1-eth1, allowing the return packets to travel from h2 to h1.

These entries effectively routed the ICMP packets between the two hosts, confirming that the network is set up correctly. The successful pings demonstrate that the OpenFlow switch is efficiently managing the traffic flow between h1 and h2 as expected.



## 9. Check the flow table of the switch again:

mininet> dpctl dump-flows

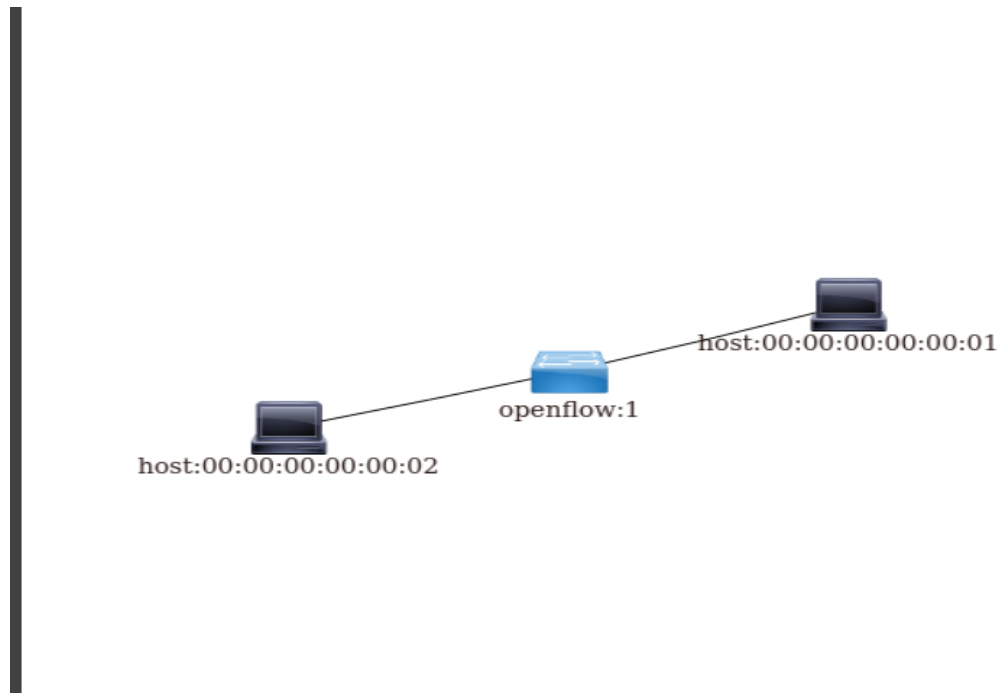
```
mininet> dpctl dump-flows
*** s1 .....
cookie=0x0, duration=66.398s, table=0, n_packets=16, n_bytes=1568, idle_timeout=60, priority=1,icmp,in_port="s1-eth2",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:02,dl_dst=00:00:00:00:01,nw_src=10.0.0.2,nw_dst=10.0.0.1,nw_tos=0,icmp_type=8,icmp_code=0 actions=output:"s1-eth1"
cookie=0x0, duration=66.394s, table=0, n_packets=16, n_bytes=1568, idle_timeout=60, priority=1,icmp,in_port="s1-eth1",vlan_tci=0x0000/0x1fff,dl_src=00:00:00:00:01,dl_dst=00:00:00:00:02,nw_src=10.0.0.1,nw_dst=10.0.0.2,nw_tos=0,icmp_type=0,icmp_code=0 actions=output:"s1-eth2"
cookie=0x0, duration=439.611s, table=0, n_packets=34, n_bytes=2572, priority=0 actions=CONTROLLER:128
mininet>
```

The successful ping tests between h1 and h2 indicate that the OpenFlow switch is properly configured. This can be directly linked to the flow table entries observed in the previous dump:

- Forwarding Traffic: There was a flow entry in the switch that managed traffic from s1-eth1 (which connects to h1) to s1-eth2 (which connects to h2), allowing packets to travel from h1 to h2.
- Handling Responses: Another flow entry was responsible for managing traffic from s1-eth2 back to s1-eth1, ensuring that the ping responses from h2 were correctly routed back to h1.

These specific flow entries were crucial in routing ICMP packets correctly between h1 and h2 and ensuring that replies were accurately delivered. The successful pings confirm that the OpenFlow switch is operating as expected and effectively handling the traffic between the hosts.

## 10. Open up OpenDaylight GUI again:



Right now this shows hosts in the topology.



## 11. Wireshark packet capture:

**\*Loopback: lo**

File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help

Apply a display filter ... <Ctrl-/> Expression... +

Packet list Narrow & Wide Case sensitive String open\_flow Find Cancel

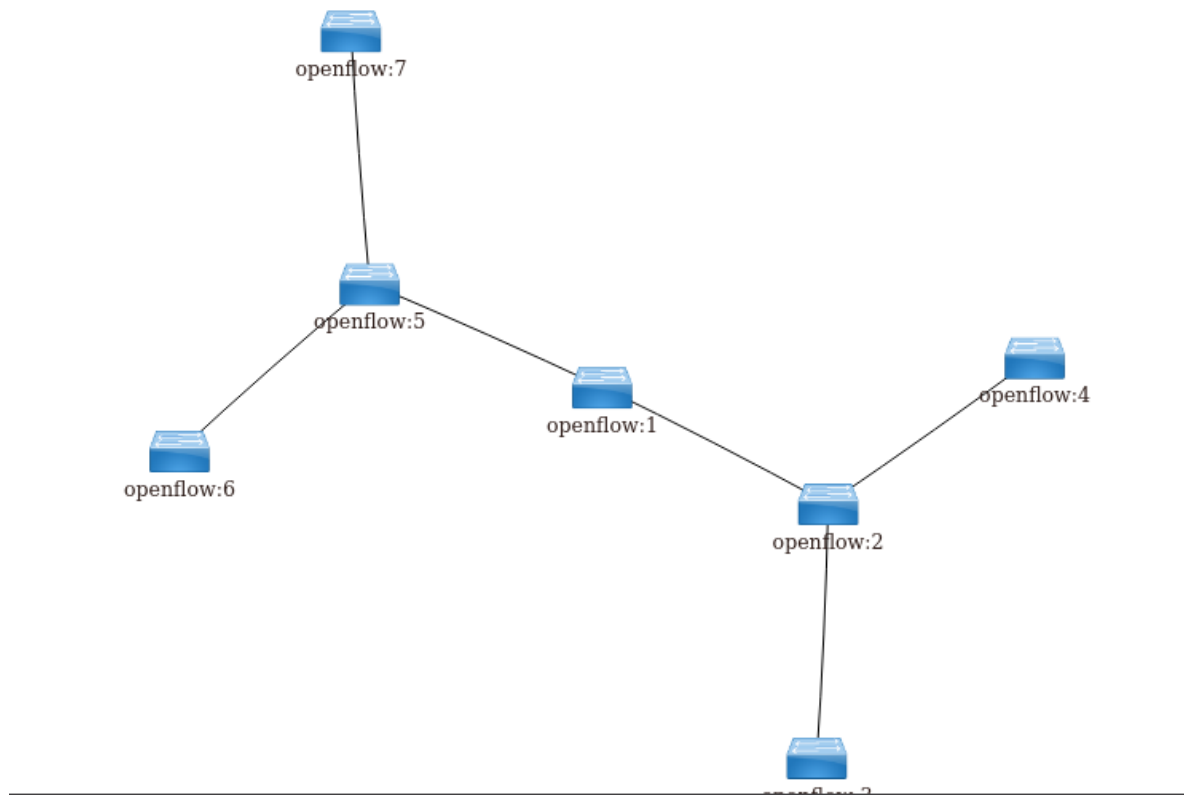
| No.  | Time            | Source     | Destination | Protocol  | Length | Info   |
|------|-----------------|------------|-------------|-----------|--------|--|
| 1070 | 1106.4915180... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 74     | Type: OFPT_ECHO_REQUEST                          |
| 1071 | 1106.4917579... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 74     | Type: OFPT_ECHO_REPLY                            |
| 1072 | 1106.4917650... | 127.0.0.1  | 127.0.0.1   | TCP       | 66     | 40274 → 6653 [ACK] Seq=4697 Ack=6563 Win=86 Len= |
| 1073 | 1111.0404800... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 178    | Type: OFPT_PACKET_IN                             |
| 1074 | 1111.0406740... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 176    | Type: OFPT_PACKET_OUT                            |
| 1075 | 1111.0406781... | 127.0.0.1  | 127.0.0.1   | TCP       | 66     | 40274 → 6653 [ACK] Seq=4809 Ack=6673 Win=86 Len= |
| 1076 | 1115.4914637... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 74     | Type: OFPT_ECHO_REQUEST                          |
| 1077 | 1115.4916437... | 127.0.0.1  | 127.0.0.1   | OpenFl... | 74     | Type: OFPT_ECHO_REPLY                            |
| 1078 | 1115.4916487... | 127.0.0.1  | 127.0.0.1   | TCP       | 66     | 40274 → 6653 [ACK] Seq=4817 Ack=6681 Win=86 Len= |
| 1079 | 1120.0191585... | 127.0.0.1  | 127.0.0.53  | DNS       | 96     | Standard query 0xdd0f A push.services.mozilla.co |
| 1080 | 1120.0192754... | 127.0.0.53 | 127.0.0.1   | DNS       | 96     | Standard query response 0xdd0f Server failure A  |

▶ Frame 64: 66 bytes on wire (528 bits), 66 bytes captured (528 bits) on interface 0  
▶ Ethernet II, Src: 00:00:00:00:00:00 (00:00:00:00:00:00), Dst: 00:00:00:00:00:00 (00:00:00:00:00:00)  
▶ Internet Protocol Version 4, Src: 127.0.0.1, Dst: 127.0.0.1  
▶ Transmission Control Protocol, Src Port: 1830, Dst Port: 42948, Seq: 1, Ack: 225, Len: 0

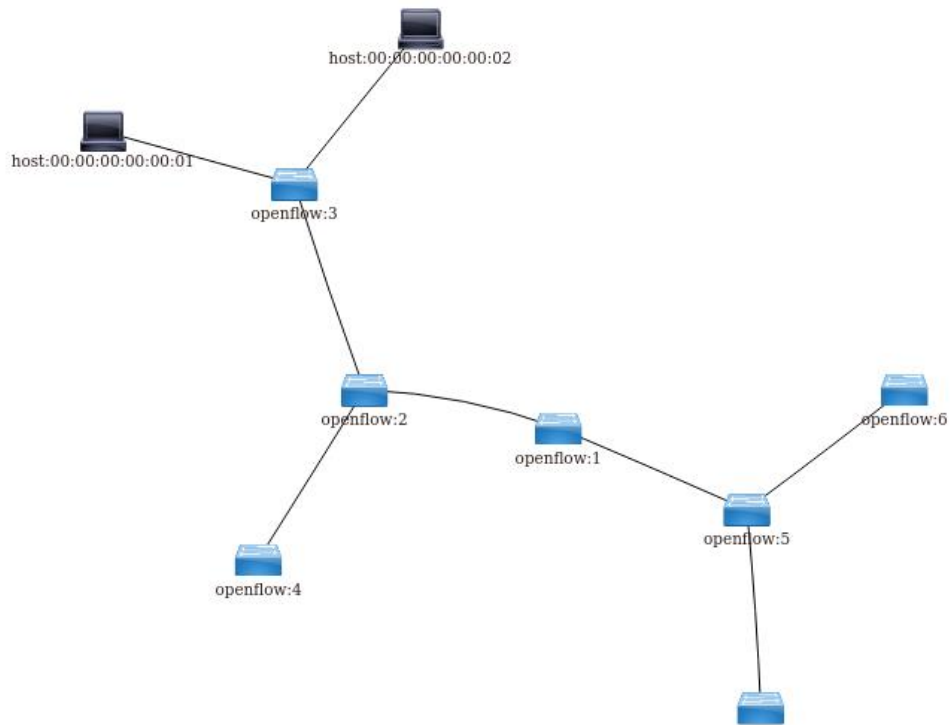
```
0000 00 00 00 00 00 00 00 00 00 00 00 08 00 45 00 .....E.
0010 00 34 b7 7c 40 00 40 06 85 45 7f 00 00 01 7f 00 -4|@.@.E.
0020 00 01 07 26 a7 c4 7d 5c fc 7e 3b 60 47 2c 80 10 ...&.:}\-;G.
0030 01 ff fe 28 00 00 01 01 08 0a 99 be 6d 2b 99 be ...(.m+
0040 6d 2b m+
```

```
sudo mn --mac --topo=tree,3 --controller=remote,ip=192.168.56.101,port=6633 --switch  
ovs,protocols=OpenFlow13
```

## Topology 02:

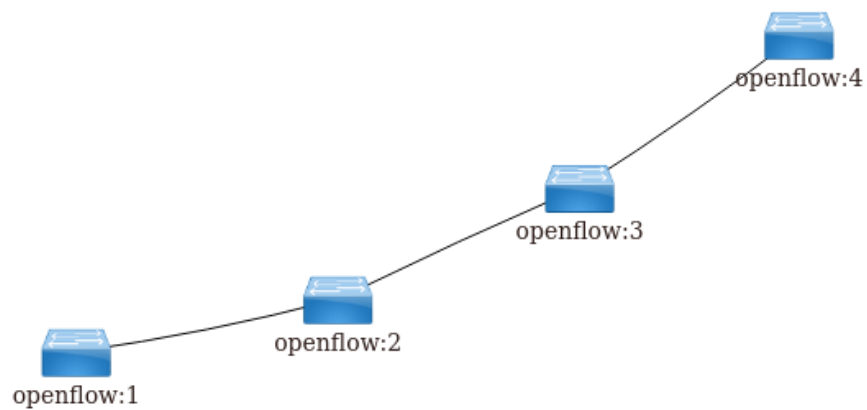


```
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.855 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.270 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.051 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.044 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.045 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.038 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.042 ms
^C
--- 10.0.0.2 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6133ms
rtt min/avg/max/mdev = 0.038/0.192/0.855/0.281 ms
mininet>
```



### Topology 03:

`sudo mn --mac --topo=linear,4 --controller=remote,ip=192.168.56.101,port=6633 --switch ovs,protocols=OpenFlow13`



```
mininet> h1 ping h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.499 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.064 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.128 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.081 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.151 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.086 ms
^C
1 --- 10.0.0.2 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5103ms
rtt min/avg/max/mdev = 0.064/0.168/0.499/0.151 ms
mininet>
```

